



深蓝学院
shenlanxueyuan.com

第七章作业讲评

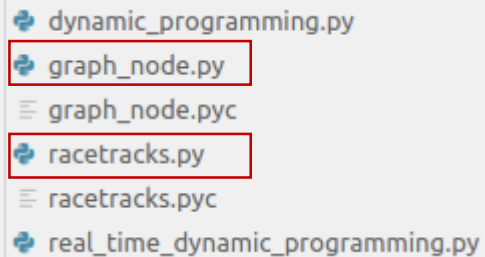


主讲人 武庆斌



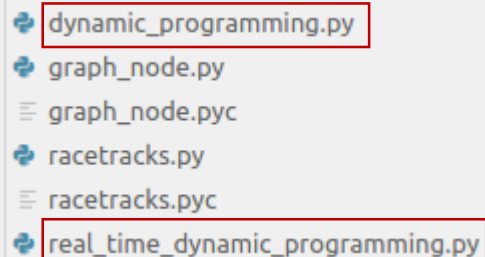
大纲

- MDP框架



A file explorer window showing a directory of files. The files listed are: dynamic_programming.py, graph_node.py, graph_node.pyc, racetracks.py, racetracks.pyc, and real_time_dynamic_programming.py. The files graph_node.py and racetracks.py are highlighted with red rectangular boxes.

- DP&RTDP流程



A file explorer window showing a directory of files. The files listed are: dynamic_programming.py, graph_node.py, graph_node.pyc, racetracks.py, racetracks.pyc, and real_time_dynamic_programming.py. The files dynamic_programming.py and real_time_dynamic_programming.py are highlighted with red rectangular boxes.

MDP formalize

Markov Decision Process (MDP)

A MDP is a 4-tuple (S, A, P, R) in learning field, or (X, U, P, L) in planning field:

- S or X is *state space*,
- A or U is *(robot) action space*,
- $P(x_{k+1}|x_k, u_k)$ is *state transition function* under probabilistic model,
 - which degenerates to a set $X_{k+1}(x_k, u_k)$ under nondeterministic model,
- $R(x_k, x_{k+1})$ is *the immediate reward*, or the negative one-step cost $-l(x_k, u_k, \theta_k)$, after transitioning from x_k to x_{k+1} due to u, θ .

MDP formalize

1.state space

① $X = \{(x, y) \mid 0 \leq x \leq 11, 0 \leq y \leq 34\}$

② $X_I = \{\text{green grids}\}$

$X_F = \{\text{yellow grids}\}$

```
START_LINE = [[0, 3], [0, 4], [0, 5], [0, 6]]
```

```
FINISH_LINE = [[34, 11], [33, 11], [32, 11]]
```

2.action space

③ $U = \{(\ddot{x}, \ddot{y}) \mid \ddot{x} \in \{0, \pm 1\}, \ddot{y} \in \{0, \pm 1\}\}$

```
ACTION_SPACE = [[1, 1], [0, 1],
```

3.state transition function

④ $\Theta = \{\theta_1, \theta_2\}$

• $\theta_1: f(\mathbf{x}_{k+1}, \mathbf{x}_k, \mathbf{u}_k) = \mathbf{x}_k \quad p_1 = 0.1$

• $\theta_2: f(\mathbf{x}_{k+1}, \mathbf{x}_k, \mathbf{u}_k) = \mathbf{x}_{k+1} \quad p_1 = 0.9$

```
def connect_to_graph(self, grid):  
    for u in ACTION_SPACE:  
        self.next_prob_9.append(self.control(u[0], u[1], grid, success=True))  
        self.next_prob_1.append(self.control(u[0], u[1], grid, success=False))
```

4.cost function

⑤ $l(\mathbf{x}_k, \mathbf{x}_k, \theta_k) = -1$

```
# TO BE IMPLEMENTED  
expected_cost_uk = 0.9 * (1+child_9.g_value) + 0.1*(1+child_1.g_value)  
value_uk.append(expected_cost_uk)
```

Initialize G values of all states to finite values;

while *not converge* **do**

for all the states x **do**

$G(x_F) = 0$;

$G_k(x_k) =$

$\min_{u_k \in U(x_k)} \{E_{\theta_k} [I(x_k, u_k, \theta_k) + G_{k+1}(x_{k+1})]\}$,

 if $x_k \neq x_F$;

end

end

DP

- 1 Initialize G values of all states to admissible values;
- 2 Follow greedy policy picking outcomes at random until goal is reached;
- 3 Backup all states visited on the way;
- 4 Reset to x_s and repeat 2-4 until all states on the current greedy policy have Bellman errors $< \Delta$, where $\Delta(x_k) = \|G(x_k) - G(x_{k+1})\|$;

RTDP

重新构建graph

```
def build_up_graph(grid, save_path):
    max_vel = 5

    # velocity dimension
    vel_list = []
    for i_vel in range(-max_vel+1, max_vel):
        for j_vel in range(-max_vel+1, max_vel):
            vel_list.append([i_vel, j_vel])

    # position dimension
    x_idx, y_idx = np.where(grid == FREE)
    coord = np.stack([x_idx, y_idx], axis=1)
    for p_idx in range(coord.shape[0]):
        pnt = coord[p_idx]
        for vel in vel_list:
            state = Node(pnt[0], pnt[1], vel[0], vel[1])
            state.connect_to_graph(grid)
            graph[state.key] = state

    for pnt in START_LINE:
        state = Node(pnt[0], pnt[1], 0, 0)
        state.connect_to_graph(grid)
        graph[state.key] = state

    for pnt in FINISH_LINE:
        state = Node(pnt[0], pnt[1], 0, 0)
        state.is_goal = True
        graph[state.key] = state

    output = open(save_path, 'wb')
    pickle.dump(graph, output)
```

DP

TIPS1: DP的g_value为默认0, RTDP的g_value需要初始化

TIPS2: 程序要跑两次, 第一次在主函数运行build_up_graph, 第二次屏蔽掉主函数的build_up_graph函数。

```
def build_up_graph(grid, save_path):
    max_vel = 5

    # velocity dimension
    vel_list = []
    for i_vel in range(-max_vel + 1, max_vel):
        for j_vel in range(-max_vel + 1, max_vel):
            vel_list.append([i_vel, j_vel])

    # position dimension
    x_idx, y_idx = np.where(grid == FREE)
    coord = np.stack([x_idx, y_idx], axis=1)
    for p_idx in range(coord.shape[0]):
        pnt = coord[p_idx]
        for vel in vel_list:
            state = Node(pnt[0], pnt[1], vel[0], vel[1])
            m_dist = np.abs(np.asarray(FINISH_LINE) - np.array([state.px, state.py]))

            # IMPORTANT-1 Heuristic Function design here
            # TO BE IMPLEMENTED

            # Note: Both the two heuristics are not the best
            # example-1 heuristic = np.linalg.norm(m_dist, axis=1) # Euclidean distance
            # example-2 heuristic = m_dist[:, 0] + m_dist[:, 1] # Mahalanobis distance
            heuristic = np.linalg.norm(m_dist, axis=1)
            state.g_value = np.min(heuristic)/3
            state.connect_to_graph(grid)
            graph[state.key] = state

    for pnt in START_LINE:
        state = Node(pnt[0], pnt[1], 0, 0)
        dist = np.abs(np.asarray(FINISH_LINE) - np.array([state.px, state.py]))
        heuristic = np.linalg.norm(dist, axis=1)
        state.g_value = np.min(heuristic)/3
        state.connect_to_graph(grid)
        graph[state.key] = state
```

RTDP

1 Initialize G values of all states to admissible values;

```
# IMPORTANT-1 Heuristic Function design here
# TO BE IMPLEMENTED

# Note: Both the two heuristics are not the best
# example-1 heuristic = np.linalg.norm(m_dist, axis=1) # Euclidean distance
# example-2 heuristic = m_dist[:, 0] + m_dist[:, 1] # Mahalanobis distance
heuristic = np.linalg.norm(m_dist, axis=1)
state.g_value = np.min(heuristic)/3
state.connect_to_graph(grid)
graph[state.key] = state

for pnt in START_LINE:
    state = Node(pnt[0], pnt[1], 0, 0)
    dist = np.abs(np.asarray(FINISH_LINE) - np.array([state.px, state.py]))
    heuristic = np.linalg.norm(dist, axis=1)
    state.g_value = np.min(heuristic)/3
    state.connect_to_graph(grid)
    graph[state.key] = state
```

value update

```
while bellman_error > 0.0001:
    itr_num += 1
    bellman_error = 0.0
    for key in graph.keys():
        state = graph[key]
        if state.is_goal:
            state.g_value = 0
        else:
            value_uk = []
            for child_idx in range(len(ACTION_SPACE)):
                child_key_9 = state.next_prob_9[child_idx]
                child_9 = graph[child_key_9]
                child_key_1 = state.next_prob_1[child_idx]
                child_1 = graph[child_key_1]
                expected_cost_uk = 0.9 * (1 + child_9.g_value) + 0.1 * (1 + child_1.g_value)
                value_uk.append(expected_cost_uk)
            current_value = min(value_uk)
            bellman_error += np.linalg.norm(state.g_value - current_value)
            state.g_value = min(value_uk)
    # end if
# end for
bellman_error_list.append(bellman_error)
print("{}th iteration: {}".format(itr_num, bellman_error))
# end while
```

DP

```
# IMPORTANT-2: implement RTDP
while bellman_error > 0.0001:
    # for i in range(500): # YOU MAY CHANGE THIS VALUE
    itr_num += 1
    bellman_error = 0.0
    rand_start = np.random.randint(low=0, high=3, size=1)[0]
    greedy_plan = greedy_policy(id=rand_start)
    for key in greedy_plan:
        state = graph[key]
        if state.is_goal:
            state.g_value = 0
        else:
            value_uk = []
            for child_idx in range(len(ACTION_SPACE)):
                child_key_9 = state.next_prob_9[child_idx]
                child_9 = graph[child_key_9]
                child_key_1 = state.next_prob_1[child_idx]
                child_1 = graph[child_key_1]

                # TO BE IMPLEMENTED
                expected_cost_uk = 0.9 * (1+child_9.g_value) + 0.1*(1+child_1.g_value)
                value_uk.append(expected_cost_uk)

                # TO BE IMPLEMENTED
                current_value = min(value_uk)
                bellman_error += np.linalg.norm(state.g_value - current_value)

                # TO BE IMPLEMENTED
                state.g_value = min(value_uk)
    # end if
# end for
bellman_error_list.append(bellman_error)
print("{}th iteration: {}".format(itr_num, bellman_error))
# end while
```

RTDP

TIPS1: DP更新graph上的所有节点，RTDP只更新已搜索到路径上的点。

value update


- 1 Initialize G values of all states to admissible values;
- 2 Follow greedy policy picking outcomes at random until goal is reached;
- 3 Backup all states visited on the way;
- 4 Reset to x_s and repeat 2-4 until all states on the current greedy policy have Bellman errors $< \Delta$, where $\Delta(x_k) = \|G(x_k) - G(x_{k+1})\|$;

```
# TO BE IMPLEMENTED
expected_cost_uk = 0.9 * (1+child_9.g_value) + 0.1*(1+child_1.g_value)
value_uk.append(expected_cost_uk)
```

```
# TO BE IMPLEMENTED
current_value = min(value_uk)
bellman_error += np.linalg.norm(state.g_value - current_value)
```


```
# TO BE IMPLEMENTED
state.g_value = min(value_uk)
```

最优路径判断



TIPS1: 由于cost function是执行一个动作加1, 所以用动作个数来判断是否达到最优, 即路径段数

TIPS2: 最优路径的判断是路径的段数, 基本是12段, 和形状没有关系



感谢各位聆听 !
Thanks for Listening

